

0. Setting

```
In [1]: # import library
import torch
import torch.nn as nn
import matplotlib.pyplot as plt
import math
from pandas import Series, DataFrame
import pandas as pd
import numpy as np

torch.__version__
```

```
Out[1]: '1.7.0+cu101'
```

1. Data

```
In [44]: from torchvision import transforms, datasets

data_path = './MNIST'

data_train = datasets.MNIST(root = data_path, train= False, download=True)
data_test = datasets.MNIST(root = data_path, train= True, download=True)

data_train_mean = data_train.data.float().mean()/255
data_train_std = data_train.data.float().std()/255

data_test_mean = data_test.data.float().mean()/255
data_test_std = data_test.data.float().std()/255

print("train data mean = {}, std = {}".format(data_train_mean, data_train_std))
print("test data mean = {}, std = {}".format(data_test_mean, data_test_std))

transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,)), # mean value = 0.1307, standard deviation value = 0.3081
])

'''
train_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((data_train_mean,), (data_train_std,)),
])

test_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((data_test_mean,), (data_test_std,)),
])
'''

data_train = datasets.MNIST(root = data_path, train= False, download=True, transform= transform)
data_test = datasets.MNIST(root = data_path, train= True, download=True, transform= transform)

print("the number of your training data (must be 10,000) = ", data_train.__len__())
print("hte number of your testing data (must be 60,000) = ", data_test.__len__())

train data mean = 0.1325146108865738, std = 0.3104802668094635
test data mean = 0.13066047430038452, std = 0.30810779333114624
the number of your training data (must be 10,000) = . 10000
hte number of your testing data (must be 60,000) = . 60000
```

2.Model

```
In [59]: def init_weights_xaiver(m):
    if type(m) == nn.Linear:
        nn.init.xavier_normal_(m.weight.data)
        m.bias.data.fill_(0)

def init_weights_he(m):
    if type(m) == nn.Linear:
        nn.init.kaiming_normal_(m.weight.data, a=0, mode='fan_in')
        m.bias.data.fill_(0)

class classification(nn.Module):

    def __init__(self):
        super(classification, self).__init__()

        # construct layers for a neural network
        self.classifier1 = nn.Sequential(
            nn.Linear(in_features=28*28, out_features=20*20),
            nn.ReLU(),
            #nn.Sigmoid(),
            #nn.Tanh(),
            #nn.LeakyReLU(),
            nn.Dropout(0.15)
        )
        self.classifier2 = nn.Sequential(
            nn.Linear(in_features=20*20, out_features=10*10),
            nn.ReLU(),
            #nn.Sigmoid(),
            #nn.Tanh(),
            #nn.LeakyReLU(),
            nn.Dropout(0.15)
        )
        self.classifier3 = nn.Sequential(
            nn.Linear(in_features=10*10, out_features=10),
```

```

        nn.LogSoftmax(dim=1)
    )

    self.classifier1.apply(init_weights_he)
    self.classifier2.apply(init_weights_he)
    self.classifier3.apply(init_weights_he)

    def forward(self, inputs):
        # [batchSize, 1, 28, 28]
        x = inputs.view(inputs.size(0), -1) # [batchSize, 28*28]
        x = self.classifier1(x) # [batchSize, 20*20]
        x = self.classifier2(x) # [batchSize, 10*10]
        out = self.classifier3(x) # [batchSize, 10]

    return out

```

3. Loss Function

```

In [4]: model = classification()
        criterion = nn.CrossEntropyLoss()
        train_y_pred = model(data_train.data.float())
        train_y = data_train.targets
        temp_loss = criterion(train_y_pred, train_y)
        print(temp_loss.data.item())

```

2.3054168224334717

4. Optimization

Define Train Function

```

In [5]: def train(model, criterion, train_loader, optimizer, batch_size):

        model.train()
        loss_sum = 0
        acc_sum = 0
        iteration = 0
        for xs, ts in iter(train_loader):

            iteration = iteration + 1
            optimizer.zero_grad()
            y_pred = model(xs)
            loss = criterion(y_pred, ts)
            loss.backward()
            optimizer.step()

            loss_sum = loss_sum + float(loss)
            zs = y_pred.max(1, keepdim=True)[1] # first column has actual prob
            acc_sum = acc_sum + zs.eq(ts.view_as(zs)).sum().item()/batch_size

        loss_avg = math.trunc(loss_sum/iteration * 100) / 100
        acc_avg = math.trunc(acc_sum/iteration * 100) / 100

        return loss_avg, acc_avg

```

Define Test Function

```

In [6]: def test(model, criterion, test_loader, batch_size):

        model.eval()
        loss_sum = 0
        acc_sum = 0
        iteration = 0
        with torch.no_grad():
            for xs, ts in iter(test_loader):
                iteration = iteration + 1
                y_pred = model(xs)
                loss_sum = loss_sum + criterion(y_pred, ts).data.item()
                zs = y_pred.max(1, keepdim=True)[1]
                acc_sum = acc_sum + zs.eq(ts.view_as(zs)).sum().item()/batch_size

        loss_avg = math.trunc(loss_sum/iteration * 100) / 100
        acc_avg = math.trunc(acc_sum/iteration * 100) / 100

        return loss_avg, acc_avg

```

Define Gradient Descent Function

```

In [7]: def gradient_descent(model, optimizer, criterion, batch_size, num_epochs):

        # batching
        train_loader = torch.utils.data.DataLoader(
            data_train,
            batch_size=batch_size,
            num_workers=2,
            shuffle=True,
            drop_last=True)

        test_loader = torch.utils.data.DataLoader(
            data_test,
            batch_size=batch_size,
            num_workers=2,
            shuffle=False,
            drop_last=True)

        # return variables
        train_loss_list, train_acc_list = [], []
        test_loss_list, test_acc_list = [], []

        # run training & testing
        for epoch in range(num_epochs + 1):

```

```

train_loss_avg, train_acc_avg = train(model, criterion, train_loader, optimizer, batch_size)
test_loss_avg, test_acc_avg = test(model, criterion, test_loader, batch_size)

# add loss and accuracy data
train_loss_list.append(train_loss_avg)
train_acc_list.append(train_acc_avg)
test_loss_list.append(test_loss_avg)
test_acc_list.append(test_acc_avg)

# print
if epoch % 10 != 0 :
    continue

print("epoch : ", epoch, " ----- ")
print("train loss : {}      accuracy = {}".format(train_loss_avg, train_acc_avg))
print("test loss : {}      accuracy = {}".format(test_loss_avg, test_acc_avg))

return train_loss_list, train_acc_list, test_loss_list, test_acc_list

```

5. Select Hyperparameter & Modify/Test Model

size of the mini-batch : 64

optimization algorithm : SGD

loss funtion : cross entropy

regularization algorithm : -

learning rate : constant

```

In [ ]: # model
        model = classification()

# mini-batch size
batch_size = 32

# num of epochs
num_epochs = 200

# learning rate
learning_rate = 0.01

# optimizer
optimizer = torch.optim.SGD(model.parameters(), lr = learning_rate)

# loss function
criterion = nn.CrossEntropyLoss()

# run
train_loss_list1, train_acc_list1, test_loss_list1, test_acc_list1 = gradient_descent(model, optimizer, criterion, batch_size, num_epochs)

```

```

epoch : 0 -----
train loss : 2.302 ..... accuracy = 0.146
test loss : 2.288 ..... accuracy = 0.136
epoch : 10 -----
train loss : 2.023 ..... accuracy = 0.761
test loss : 2.018 ..... accuracy = 0.753
epoch : 20 -----
train loss : 1.843 ..... accuracy = 0.821
test loss : 1.846 ..... accuracy = 0.808
epoch : 30 -----
train loss : 1.761 ..... accuracy = 0.845
test loss : 1.769 ..... accuracy = 0.828
epoch : 40 -----
train loss : 1.713 ..... accuracy = 0.86
test loss : 1.722 ..... accuracy = 0.845
epoch : 50 -----
train loss : 1.681 ..... accuracy = 0.872
test loss : 1.691 ..... accuracy = 0.858
epoch : 60 -----
train loss : 1.657 ..... accuracy = 0.882
test loss : 1.669 ..... accuracy = 0.866
epoch : 70 -----
train loss : 1.639 ..... accuracy = 0.888
test loss : 1.652 ..... accuracy = 0.871
epoch : 80 -----
train loss : 1.626 ..... accuracy = 0.893
test loss : 1.64 ..... accuracy = 0.876
epoch : 90 -----
train loss : 1.615 ..... accuracy = 0.897
test loss : 1.629 ..... accuracy = 0.879
epoch : 100 -----
train loss : 1.606 ..... accuracy = 0.9
test loss : 1.621 ..... accuracy = 0.882
epoch : 110 -----
train loss : 1.599 ..... accuracy = 0.905
test loss : 1.614 ..... accuracy = 0.885
epoch : 120 -----
train loss : 1.593 ..... accuracy = 0.907
test loss : 1.608 ..... accuracy = 0.887
epoch : 130 -----
train loss : 1.587 ..... accuracy = 0.91
test loss : 1.603 ..... accuracy = 0.889
epoch : 140 -----
train loss : 1.581 ..... accuracy = 0.913
test loss : 1.598 ..... accuracy = 0.891
epoch : 150 -----
train loss : 1.576 ..... accuracy = 0.915
test loss : 1.593 ..... accuracy = 0.894
epoch : 160 -----
train loss : 1.571 ..... accuracy = 0.918
test loss : 1.589 ..... accuracy = 0.895
epoch : 170 -----
train loss : 1.567 ..... accuracy = 0.919
test loss : 1.585 ..... accuracy = 0.898
epoch : 180 -----
train loss : 1.563 ..... accuracy = 0.921
test loss : 1.582 ..... accuracy = 0.899
epoch : 190 -----
train loss : 1.56 ..... accuracy = 0.923

```

```

test loss : 1.579 ..... accuracy = 0.9
epoch : 200 .....
train loss : 1.557 ..... accuracy = 0.925
test loss : 1.576 ..... accuracy = 0.902

```

```

In [8]: def gradient_descent_with_scheduler(scheduler, model, optimizer, criterion, batch_size, num_epochs):

    # batching
    train_loader = torch.utils.data.DataLoader(
        data_train,
        batch_size=batch_size,
        num_workers=2,
        shuffle=True,
        drop_last=True)

    test_loader = torch.utils.data.DataLoader(
        data_test,
        batch_size=batch_size,
        num_workers=2,
        shuffle=False,
        drop_last=True)

    # return variables
    train_loss_list, train_acc_list = [], []
    test_loss_list, test_acc_list = [], []

    # run training & testing
    for epoch in range(num_epochs + 1):

        train_loss_avg, train_acc_avg = train(model, criterion, train_loader, optimizer, batch_size)
        test_loss_avg, test_acc_avg = test(model, criterion, test_loader, batch_size)
        scheduler.step(train_loss_avg)

        # add loss and accuracy data
        train_loss_list.append(train_loss_avg)
        train_acc_list.append(train_acc_avg)
        test_loss_list.append(test_loss_avg)
        test_acc_list.append(test_acc_avg)

        # print
        if epoch % 10 != 0 :
            continue

        print("epoch : ", epoch, " ----- ")
        print("train loss : {}      accuracy = {}".format(train_loss_avg, train_acc_avg))
        print("test loss : {}      accuracy = {}".format(test_loss_avg, test_acc_avg))

    return train_loss_list, train_acc_list, test_loss_list, test_acc_list

```

```

In [61]: # model
model = classification()

# mini-batch size
batch_size = 32

# num of epochs
num_epochs = 100

# learning rate
learning_rate = 0.01

# optimizer
optimizer = torch.optim.SGD(model.parameters(), lr = learning_rate, weight_decay=0.0001)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'min', factor=0.3, verbose=True)

# loss function
criterion = nn.CrossEntropyLoss()

# run
train_loss_list8, train_acc_list8, test_loss_list8, test_acc_list8 = gradient_descent_with_scheduler(scheduler, model, optimizer, criterion, batch_size, num_epochs)

epoch : 0 .....
train loss : 0.85 ..... accuracy = 0.72
test loss : 0.42 ..... accuracy = 0.88
epoch : 10 .....
train loss : 0.13 ..... accuracy = 0.95
test loss : 0.2 ..... accuracy = 0.94
epoch : 20 .....
train loss : 0.06 ..... accuracy = 0.97
test loss : 0.18 ..... accuracy = 0.95
epoch : 30 .....
train loss : 0.03 ..... accuracy = 0.98
test loss : 0.17 ..... accuracy = 0.95
epoch : 40 .....
train loss : 0.03 ..... accuracy = 0.99
test loss : 0.17 ..... accuracy = 0.95
Epoch : 51: reducing learning rate of group 0 to 3.0000e-03.
epoch : 50 .....
train loss : 0.02 ..... accuracy = 0.99
test loss : 0.17 ..... accuracy = 0.95
epoch : 60 .....
train loss : 0.01 ..... accuracy = 0.99
test loss : 0.18 ..... accuracy = 0.95
Epoch : 63: reducing learning rate of group 0 to 9.0000e-04.
epoch : 70 .....
train loss : 0.01 ..... accuracy = 0.99
test loss : 0.18 ..... accuracy = 0.96
Epoch : 74: reducing learning rate of group 0 to 2.7000e-04.
epoch : 80 .....
train loss : 0.01 ..... accuracy = 0.99
test loss : 0.18 ..... accuracy = 0.96
Epoch : 85: reducing learning rate of group 0 to 8.1000e-05.
epoch : 90 .....
train loss : 0.01 ..... accuracy = 0.99
test loss : 0.17 ..... accuracy = 0.96
Epoch : 96: reducing learning rate of group 0 to 2.4300e-05.
epoch : 100 .....

```

```
train loss : 0.01 ..... accuracy = 0.99
test loss  : 0.17 ..... accuracy = 0.96
```

```
In [27]: # model
model = classification()

# mini-batch size
batch_size = 64

# num of epochs
num_epochs = 100

# learning rate
learning_rate = 0.01

# optimizer
optimizer = torch.optim.SGD(model.parameters(), lr = learning_rate, weight_decay=0.0001)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'min', factor=0.5, patience = 5, verbose=True)

# loss function
criterion = nn.CrossEntropyLoss()

# run
train_loss_list7, train_acc_list7, test_loss_list7, test_acc_list7 = gradient_descent_with_scheduler(scheduler, model, optimizer, criterion, batch_size, num_epochs)


epoch : 0 .....
train loss : 1.84 ..... accuracy = 0.36
test loss  : 1.25 ..... accuracy = 0.62
epoch : 10 .....
train loss : 0.24 ..... accuracy = 0.92
test loss  : 0.25 ..... accuracy = 0.93
epoch : 20 .....
train loss : 0.15 ..... accuracy = 0.95
test loss  : 0.2 ..... accuracy = 0.94
epoch : 30 .....
train loss : 0.11 ..... accuracy = 0.96
test loss  : 0.18 ..... accuracy = 0.95
epoch : 40 .....
train loss : 0.08 ..... accuracy = 0.97
test loss  : 0.18 ..... accuracy = 0.95
epoch : 50 .....
train loss : 0.06 ..... accuracy = 0.98
test loss  : 0.17 ..... accuracy = 0.95
epoch : 60 .....
train loss : 0.05 ..... accuracy = 0.98
test loss  : 0.17 ..... accuracy = 0.95
Epoch ... 63: reducing learning rate of group 0 to 5.0000e-03.
epoch : 70 .....
train loss : 0.04 ..... accuracy = 0.98
test loss  : 0.17 ..... accuracy = 0.95
Epoch ... 72: reducing learning rate of group 0 to 2.5000e-03.
Epoch ... 78: reducing learning rate of group 0 to 1.2500e-03.
epoch : 80 .....
train loss : 0.04 ..... accuracy = 0.98
test loss  : 0.17 ..... accuracy = 0.95
Epoch ... 84: reducing learning rate of group 0 to 6.2500e-04.
Epoch ... 90: reducing learning rate of group 0 to 3.1250e-04.
epoch : 90 .....
train loss : 0.04 ..... accuracy = 0.98
test loss  : 0.17 ..... accuracy = 0.95
Epoch ... 96: reducing learning rate of group 0 to 1.5625e-04.
epoch : 100 .....
train loss : 0.04 ..... accuracy = 0.98
test loss  : 0.17 ..... accuracy = 0.95
```

In []:

6. Output


1. Plot the training and testing losses over epochs [2pt]

```
In [62]: plt.title("Loss")
plt.plot(train_loss_list8, c = 'red', label = 'train loss')
plt.plot(test_loss_list8, c = 'blue', label = 'test loss')
plt.legend(loc = 'lower right')
plt.show()
```



1. Plot the training and testing accuracies over epochs [2pt]

```
In [63]: plt.title("Accuracy")
plt.plot(train_acc_list8, c = 'red', label = 'train accuracy')
plt.plot(test_acc_list8, c = 'blue', label = 'test accuracy')
plt.legend(loc = 'lower right')
plt.show()
```



1. Print the final training and testing losses at convergence [2pt]

```
In [66]: data1 = {'': [train_loss_list8[-1], test_loss_list8[-1]]}
index1 = ['training', 'testing']
frame1 = DataFrame(data1, index = index1)
frame1.columns.name = 'loss'
frame1
```

```
Out[66]:
```

loss	
training	0.01
testing	0.17

1. Print the final training and testing accuracies at convergence [20pt]

```
In [73]: data2 = {'': [train_acc_list8[-1], test_acc_list8[-1]]}
index2 = ['training', 'testing']
frame2 = DataFrame(data2, index = index2)
frame2.columns.name = 'accuracy'
frame2
```

```
Out[73]:
```

accuracy	
training	0.99
testing	0.96