

▼ 0. Setting

```
# import library
import torch
import torch.nn as nn
import matplotlib.pyplot as plt
import math
from pandas import Series, DataFrame
import pandas as pd
import numpy as np
import torch.nn.functional as F
import math

torch.__version__

'1.7.0+cu101'

# using gpu
use_cuda = True
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(torch.cuda.is_available())
print(device)

True
cuda
```

▼ 1. Data

```
from torchvision import transforms, datasets

transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,)), # mean value = 0.1307, standard deviation v:
])

data_path = './MNIST'

data_test = datasets.MNIST(root = data_path, train= True, download=True, transform= trans
data_train = datasets.MNIST(root = data_path, train= False, download=True, transform= tran

print("the number of your training data (must be 10,000) = ", data_train.__len__())
print("hte number of your testing data (must be 60,000) = ", data_test.__len__())
```

```

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./MNIST/MNIST/raw/
9920512/? [00:03<00:00, 2643255.26it/s]
Extracting ./MNIST/MNIST/raw/train-images-idx3-ubyte.gz to ./MNIST/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./MNIST/MNIST/raw/
32768/? [00:01<00:00, 16690.68it/s]
Extracting ./MNIST/MNIST/raw/train-labels-idx1-ubyte.gz to ./MNIST/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ./MNIST/MNIST/raw/t
1654784/? [00:01<00:00, 1049065.26it/s]
Extracting ./MNIST/MNIST/raw/t10k-images-idx3-ubyte.gz to ./MNIST/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ./MNIST/MNIST/raw/t
0% 0/4542 [00:00<?, ?it/s]
Extracting ./MNIST/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./MNIST/MNIST/raw
Processing...
Done!
the number of your training data (must be 10,000) = 10000
hte number of your testing data (must be 60,000) = 60000
/usr/local/lib/python3.6/dist-packages/torchvision/datasets/mnist.py:480: UserWarning: The gi
return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)

```

▼ 2. Model

```

class MyModel(nn.Module):

    def __init__(self, num_classes=10, size_kernel=5):

        super(MyModel, self).__init__()

        # *****
        # input parameter
        #
        # data size:
        # mnist : 28 * 28
        #
        # the tensor given to the model should be of shape [batch_size, 1, height, width]
        # because first convolution has in_channels = 1
        # *****
        self.number_class = num_classes
        self.size_kernel = size_kernel

        # *****
        # feature layer
        # *****
        self.conv1 = nn.Conv2d(1, 20, kernel_size=size_kernel, stride=1, padding=i
        self.conv2 = nn.Conv2d(20, 50, kernel_size=size_kernel, stride=1, padding=

```

```

self.conv_layer1 = nn.Sequential(self.conv1, nn.MaxPool2d(kernel_size=2), nn.ReLU)
self.conv_layer2 = nn.Sequential(self.conv2, nn.MaxPool2d(kernel_size=2), nn.ReLU)

```

```

self.feature = nn.Sequential(self.conv_layer1, self.conv_layer2)

```

```

# *****

```

```

# classifier layer

```

```

# *****

```

```

self.fc1 = nn.Linear(50*7*7, 50, bias=True)

```

```

self.fc2 = nn.Linear(50, num_classes, bias=True)

```

```

self.fc_layer1 = nn.Sequential(self.fc1, nn.ReLU(True))

```

```

self.fc_layer2 = nn.Sequential(self.fc2, nn.LogSoftmax(dim=1))

```

```

self.classifier = nn.Sequential(self.fc_layer1, self.fc_layer2)

```

```

# *****

```

```

# dropout

```

```

# *****

```

```

self.dropout1 = nn.Dropout(0.25)

```

```

self.dropout2 = nn.Dropout(0.5)

```

```

self._initialize_weight()

```

```

def _initialize_weight(self):

```

```

    for m in self.modules():

```

```

        if isinstance(m, nn.Conv2d):

```

```

            #nn.init.xavier_uniform_(m.weight, gain=math.sqrt(2))

```

```

            nn.init.kaiming_normal_(m.weight.data, a=0, mode='fan_in')

```

```

            if m.bias is not None:

```

```

                m.bias.data.zero_()

```

```

        elif isinstance(m, nn.Linear):

```

```

            #nn.init.xavier_uniform_(m.weight, gain=math.sqrt(2))

```

```

            nn.init.kaiming_normal_(m.weight.data, a=0, mode='fan_in')

```

```

            if m.bias is not None:

```

```

                m.bias.data.zero_()

```

```

def forward(self, x):

```

```

    x = x.to(device)

```

```

    x = self.feature(x)

```

```

    x = self.dropout1(x)

```

```

    x = x.view(x.size(0), -1)

```

```

    x = self.dropout2(x)

```

```
x = self.dropout2(x)
x = self.classifier(x)
```

```
return x
```

▼ 3. Loss Function

```
model = MyModel(10, 5).to(device)
criterion = nn.CrossEntropyLoss()
train_y_pred = model.forward(data_train.data.unsqueeze(dim=1).float())
train_y = data_train.targets.to(device)
temp_loss = criterion(train_y_pred, train_y)
print(temp_loss.data.item())
```

```
776.919677734375
```

▼ 4. Optimization

Define Train Function

```
def train(model, criterion, train_loader, optimizer, batch_size):
```

```
    model.train()
    loss_sum = 0
    acc_sum = 0
    iteration = 0
    for xs, ts in iter(train_loader):
```

```
        iteration = iteration + 1
        optimizer.zero_grad()
        y_pred = model(xs)
        ts = ts.to(device)
        loss = criterion(y_pred, ts)
        loss.backward()
        optimizer.step()
```

```
        loss_sum = loss_sum + float(loss)
        zs = y_pred.max(1, keepdim=True)[1] # first column has actual prob
        acc_sum = acc_sum + zs.eq(ts.view_as(zs)).sum().item()/batch_size
```

```
    loss_avg = round(loss_sum/iteration, 5)
    acc_avg = round(acc_sum/iteration, 5)
```

```
    return loss_avg, acc_avg
```

Define Test Function

```
def test(model, criterion, test_loader, batch_size):
    model.eval()
    loss_sum = 0
    acc_sum = 0
    iteration = 0
    with torch.no_grad():
        for xs, ts in iter(test_loader):
            iteration = iteration + 1
            ts = ts.to(device)
            y_pred = model(xs)
            loss_sum = loss_sum + criterion(y_pred, ts).data.item()
            zs = y_pred.max(1, keepdim=True)[1]
            acc_sum = acc_sum + zs.eq(ts.view_as(zs)).sum().item()/batch_size

    loss_avg = round(loss_sum/iteration, 5)
    acc_avg = round(acc_sum/iteration, 5)

    return loss_avg, acc_avg
```

Define Gradient Descent Function

```
def gradient_descent(model, optimizer, criterion, batch_size, num_epochs):

    # batching
    train_loader = torch.utils.data.DataLoader(
        data_train,
        batch_size=batch_size,
        num_workers=4,
        shuffle=True,
        drop_last=True)

    test_loader = torch.utils.data.DataLoader(
        data_test,
        batch_size=batch_size,
        num_workers=4,
        shuffle=False,
        drop_last=True)

    # return variables
    train_loss_list, train_acc_list = [], []
    test_loss_list, test_acc_list = [], []

    # run training & testing
    for epoch in range(num_epochs + 1):

        train_loss_avg, train_acc_avg = train(model, criterion, train_loader, optimizer, batch_size)
        test_loss_avg, test_acc_avg = test(model, criterion, test_loader, batch_size)

        # add loss and accuracy data
        train_loss_list.append(train_loss_avg)
```

```

train_loss_list.append(train_loss_avg)
train_acc_list.append(train_acc_avg)
test_loss_list.append(test_loss_avg)
test_acc_list.append(test_acc_avg)

# print
if epoch % 10 != 0 :
    continue

print("epoch : ", epoch, " ----- ")
print("train loss : {}          accuracy = {}".format(train_loss_avg, train_acc_avg))
print("test loss : {}          accuracy = {}".format(test_loss_avg, test_acc_avg))

```

```

return train_loss_list, train_acc_list, test_loss_list, test_acc_list

```

```

def gradient_descent_with_scheduler(scheduler, model, optimizer, criterion, batch_size, num

```

```

# batching
train_loader = torch.utils.data.DataLoader(
    data_train,
    batch_size=batch_size,
    num_workers=4,
    shuffle=True,
    drop_last=True)

```

```

test_loader = torch.utils.data.DataLoader(
    data_test,
    batch_size=batch_size,
    num_workers=4,
    shuffle=False,
    drop_last=True)

```

```

# return variables
#train_loss_list, train_acc_list = [], []
#test_loss_list, test_acc_list = [], []

```

```

# run training & testing
for epoch in range(num_epochs + 1):

```

```

    train_loss_avg, train_acc_avg = train(model, criterion, train_loader, optimizer, batch_
    test_loss_avg, test_acc_avg = test(model, criterion, test_loader, batch_size)
    scheduler.step(train_loss_avg)

```

```

# add loss and accuracy data
train_loss_list.append(train_loss_avg)
train_acc_list.append(train_acc_avg)
test_loss_list.append(test_loss_avg)
test_acc_list.append(test_acc_avg)

```

```

# print
if epoch % 10 != 0 :

```

```
continue
```

```
print("epoch : ", epoch, " ----- ")
print("train loss : {}          accuracy = {}".format(train_loss_avg, train_acc_avg))
print("test loss : {}          accuracy = {}".format(test_loss_avg, test_acc_avg))
```

▼ 5. Plot Function

```
def plot_loss(train_loss_list, test_loss_list):
    plt.title("Loss")
    plt.plot(train_loss_list, c = 'red', label = 'train loss')
    plt.plot(test_loss_list, c = 'blue', label = 'test loss')
    plt.legend(loc = 'upper right')
    plt.show()

def plot_accuracy(train_acc_list, test_acc_list):
    plt.title("Accuracy")
    plt.plot(train_acc_list, c = 'red', label = 'train accuracy')
    plt.plot(test_acc_list, c = 'blue', label = 'test accuracy')
    plt.legend(loc = 'lower right')
    plt.show()
```

▼ 6. Run

```
# model
num_classes=10
size_kernel=5
model1 = MyModel(num_classes, size_kernel).to(device)

# mini-batch size
batch_size = 32

# num of epochs
num_epochs = 50

# learning rate
learning_rate = 0.01

# optimizer
optimizer = torch.optim.SGD(model1.parameters(), lr = learning_rate, weight_decay=0.0001)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'min', factor=0.5, patience=10)


# loss function
criterion = nn.CrossEntropyLoss()

# run
```

```
train_loss_list1, train_acc_list1, test_loss_list1, test_acc_list1 = [], [], [], []
gradient_descent_with_scheduler(scheduler, model1, optimizer, criterion, batch_size, num_epochs)
```

```
# plot
plot_loss(train_loss_list1, test_loss_list1)
```

```
epoch : 0 -----
train loss : 0.62      accuracy = 0.8
test loss : 0.24      accuracy = 0.92
epoch : 10 -----
train loss : 0.03      accuracy = 0.98
test loss : 0.09      accuracy = 0.97
epoch : 20 -----
train loss : 0.01      accuracy = 0.99
test loss : 0.07      accuracy = 0.97
Epoch 24: reducing learning rate of group 0 to 5.0000e-03.
epoch : 30 -----
train loss : 0.0       accuracy = 0.99
test loss : 0.07      accuracy = 0.97
Epoch 33: reducing learning rate of group 0 to 2.5000e-03.
Epoch 39: reducing learning rate of group 0 to 1.2500e-03.
epoch : 40 -----
train loss : 0.0       accuracy = 0.99
test loss : 0.07      accuracy = 0.97
Epoch 45: reducing learning rate of group 0 to 6.2500e-04.
Epoch 51: reducing learning rate of group 0 to 3.1250e-04.
epoch : 50 -----
train loss : 0.0       accuracy = 0.99
test loss : 0.07      accuracy = 0.97
```



```
# model
num_classes=10
size_kernel=5
model2 = MyModel(num_classes, size_kernel).to(device)
```

```
# mini-batch size
batch_size = 32
```

```
# num of epochs
num_epochs = 100
```

```
# learning rate
learning_rate = 0.01
```



```
learning_rate = 0.01
```

```
# optimizer
```

```
optimizer = torch.optim.SGD(model2.parameters(), lr = learning_rate, weight_decay=0.0001)
```

```
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'min', factor=0.5, patience=10)
```

```
# loss function
```

```
criterion = nn.CrossEntropyLoss()
```

```
# run
```

```
train_loss_list2, train_acc_list2, test_loss_list2, test_acc_list2 = [], [], [], []
```

```
gradient_descent_with_scheduler(scheduler, model2, optimizer, criterion, batch_size, num_epochs)
```

```
# plot
```

```
plot_loss(train_loss_list2, test_loss_list2)
```

```

epoch : 0 -----
train loss : 0.84726      accuracy = 0.72436
test loss : 0.26487      accuracy = 0.91923
epoch : 10 -----
train loss : 0.0737      accuracy = 0.97646
test loss : 0.0795      accuracy = 0.97597
epoch : 20 -----
train loss : 0.04198      accuracy = 0.98718
test loss : 0.06594      accuracy = 0.9801
Epoch 30: reducing learning rate of group 0 to 5.0000e-03.
epoch : 30 -----

```

```

# model
num_classes=10
size_kernel=7
model3 = MyModel(num_classes, size_kernel).to(device)

# mini-batch size
batch_size = 32

# num of epochs
num_epochs = 100

# learning rate
learning_rate = 0.01

# optimizer
optimizer = torch.optim.SGD(model3.parameters(), lr = learning_rate, weight_decay=0.0001)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'min', factor=0.5, patience=10)

# loss function
criterion = nn.CrossEntropyLoss()

# run
train_loss_list3, train_acc_list3, test_loss_list3, test_acc_list3 = [], [], [], []
gradient_descent_with_scheduler(scheduler, model3, optimizer, criterion, batch_size, num_epochs)

# plot
plot_loss(train_loss_list3, test_loss_list3)

```

```

epoch : 0 -----
train loss : 0.87445      accuracy = 0.70833
test loss : 0.24402      accuracy = 0.93017
epoch : 10 -----
train loss : 0.07214      accuracy = 0.97486
test loss : 0.0816       accuracy = 0.97485
epoch : 20 -----
train loss : 0.03971      accuracy = 0.98688
test loss : 0.07057      accuracy = 0.97907
epoch : 30 -----
train loss : 0.02691      accuracy = 0.99219
test loss : 0.07332      accuracy = 0.97922
epoch : 40 -----
train loss : 0.01742      accuracy = 0.99439
test loss : 0.07023      accuracy = 0.98115
epoch : 50 -----
train loss : 0.0164       accuracy = 0.99479
test loss : 0.0702       accuracy = 0.98198
epoch : 60 -----
train loss : 0.01049      accuracy = 0.997
test loss : 0.06953      accuracy = 0.9824
Epoch 66: reducing learning rate of group 0 to 5.0000e-03.
epoch : 70 -----
train loss : 0.00873      accuracy = 0.9973
test loss : 0.06852      accuracy = 0.983
epoch : 80 -----
train loss : 0.00945      accuracy = 0.9974
test loss : 0.06948      accuracy = 0.9827
Epoch 85: reducing learning rate of group 0 to 2.5000e-03.
epoch : 90 -----
train loss : 0.00486      accuracy = 0.999
test loss : 0.06863      accuracy = 0.9832
Epoch 97: reducing learning rate of group 0 to 1.2500e-03.
epoch : 100 -----

```

```

# num of epochs
num_epochs = 10

```

```

gradient_descent_with_scheduler(scheduler, model3, optimizer, criterion, batch_size, num_epochs)

```

```

epoch : 0 -----
train loss : 0.00588      accuracy = 0.9986
test loss : 0.06838      accuracy = 0.98342
Epoch 103: reducing learning rate of group 0 to 6.2500e-04.
Epoch 109: reducing learning rate of group 0 to 3.1250e-04.
epoch : 10 -----
train loss : 0.00612      accuracy = 0.9981
test loss : 0.06843      accuracy = 0.98352

```



```

# num of epochs
num_epochs = 50

```

```

gradient_descent_with_scheduler(scheduler, model3, optimizer, criterion, batch_size, num_epochs)

```

```

epoch : 0 -----
train loss : 0.00628      accuracy = 0.9982
test loss : 0.06825      accuracy = 0.9835

```

```

Epoch 115: reducing learning rate of group 0 to 1.5625e-04.
Epoch 121: reducing learning rate of group 0 to 7.8125e-05.
epoch : 10 -----
train loss : 0.0081      accuracy = 0.997
test loss : 0.06801     accuracy = 0.98352
Epoch 127: reducing learning rate of group 0 to 3.9063e-05.
Epoch 133: reducing learning rate of group 0 to 1.9531e-05.
epoch : 20 -----
train loss : 0.00721     accuracy = 0.9978
test loss : 0.06792     accuracy = 0.98347
Epoch 139: reducing learning rate of group 0 to 9.7656e-06.
epoch : 30 -----
train loss : 0.00603     accuracy = 0.9984
test loss : 0.06793     accuracy = 0.9835
Epoch 145: reducing learning rate of group 0 to 4.8828e-06.
Epoch 151: reducing learning rate of group 0 to 2.4414e-06.
epoch : 40 -----
train loss : 0.00484     accuracy = 0.999
test loss : 0.06793     accuracy = 0.9835
Epoch 159: reducing learning rate of group 0 to 1.2207e-06.
epoch : 50 -----
train loss : 0.00594     accuracy = 0.9984
test loss : 0.06793     accuracy = 0.98348

```

```

length = len(test_acc_list3)-10
print(length)


```

153

▼ 7. Output


1. Plot the training and testing losses over epochs [2pt]

```
plot_loss(train_loss_list3[0:length], test_loss_list3[0:length])
```



2. Plot the training and testing accuracies over epochs [2pt]

```
plot_accuracy(train_acc_list3[0:length], test_acc_list3[0:length])
```



3. Print the final training and testing losses at convergence [2pt]

```
data1 = {'': [train_loss_list3[length-1], test_loss_list3[length-1]]}
index1 = ['training', 'testing']
frame1 = DataFrame(data1, index = index1)
frame1.columns.name = 'loss'
frame1
```

	loss
training	0.00484
testing	0.06793

4. Print the final training and testing accuracies at convergence [20pt]

```
data2 = {'': [str(train_acc_list3[length-1]) + '00', str(test_acc_list3[length-1]) + '0']}
index2 = ['training', 'testing']
frame2 = DataFrame(data2, index = index2)
frame2.columns.name = 'accuracy'
frame2
```

	accuracy
training	0.99900
testing	0.98350

5. Print the testing accuracies within the last 10 epochs [5pt]

```
for i in range(10):
```

```
for i in range(10):  
    idx = length - 10 + i  
    print("[epoch = {0}] {1:0.5f}".format(idx, test_acc_list3[idx]))
```

```
[epoch = 143] 0.98350  
[epoch = 144] 0.98350  
[epoch = 145] 0.98350  
[epoch = 146] 0.98350  
[epoch = 147] 0.98350  
[epoch = 148] 0.98350  
[epoch = 149] 0.98350  
[epoch = 150] 0.98350  
[epoch = 151] 0.98350  
[epoch = 152] 0.98350
```